

Lecture 11

Lecture 11

Density estimation

Intro

Parametric methods: generative models

Parametric methods: estimation

MLE for multinomials

Nonparametric methods

Kernel

Bandwidth

Naive Bayes

Estimation

A naive assumption

Example: Discrete features

Discrete features: More formally

Continuous features

Naive Bayes: Prediction

What is Naive Bayes learning? Connection to logistic regression

Multi-armed bandits (Reinforcement Learning)

Decision making

Multiarmed bandits: Motivation and Applications

Formal Setup

Evaluating performance

Environments

Empirical means

Exploitation only

Explore-then-Exploit

ϵ -Greedy: A slightly better algorithm

UCB: More adaptive exploration

Limitations of multi-armed bandits

Responsible ML

ML models can be very sensitive to changes in the data distribution

ML models can be biased against certain sub-populations

Conclusion

Density estimation

Intro

With clustering using GMMs, our high-level goal was the following:

Given a training set x_1, \dots, x_n , estimate a density function p that could have generated this dataset (via $x_i \sim^{i.i.d.} p$).

This is a special case of the general problem of density estimation, an important unsupervised learning problem.

Density estimation is useful for many downstream applications

- we have seen clustering already, will see more today
- these applications also provide a way to measure quality of the density estimator

Parametric methods: generative models

Parametric estimation assumes a generative model parametrized by θ :

$$p(x) = p(x; \theta)$$

Examples:

- GMM: $p(x; \theta) = \sum_{j=1}^k \pi_j N(x | \mu_j, \Sigma_j)$ where $\theta = \{\pi_j, \mu_j, \Sigma_j\}$
- Multinomial: a discrete variable with values in $\{1, 2, \dots, k\}$ s.t.

$$p(x = j; \theta) = \theta_j$$

where θ is a distribution over the k elements.

Size of θ is independent of the size of the training set, so it's parametric.

Parametric methods: estimation

As usual, if i.i.d, we can apply MLE to learn the parameters θ :

$$\arg \max_{\theta} \sum_{i=1}^n \ln p(x_i; \theta)$$

For some cases this is intractable and we can use algorithms such as EM to approximately solve the MLE problem (e.g. GMMS).

For some other cases this admits a simple closed-form solution (e.g. multinomial).

MLE for multinomials

The log-likelihood is

$$\begin{aligned} \sum_{i=1}^n \ln p(x = x_i; \theta) &= \sum_{i=1}^n \ln \theta_{x_i} = \sum_{i=1}^n \sum_{j=1}^k \mathbb{1}(x_i = j) \ln \theta_j \\ &= \sum_{j=1}^k \sum_{i:x_i=j} \ln \theta_j = \sum_{j=1}^k z_j \ln \theta_j \end{aligned}$$

where $z_j = |\{i : x_i = j\}|$ is the number of examples with value j .

The solution is simply

$$\theta_j = \frac{z_j}{n} \propto z_j$$

i.e. the fraction of examples with value j . (HW4 Q2.1)

Nonparametric methods

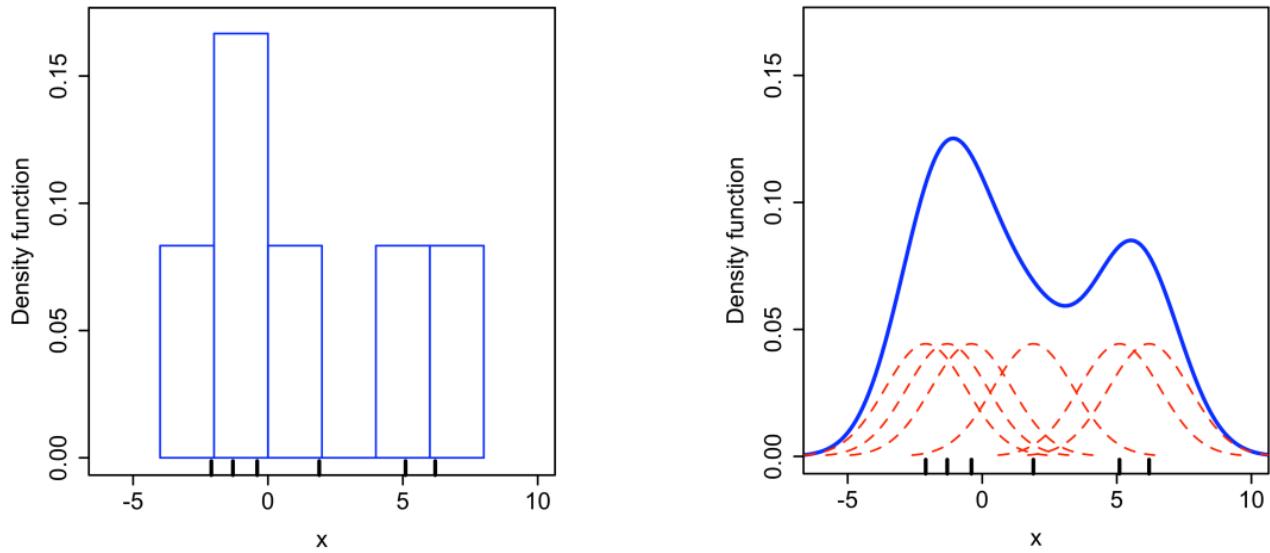
Can we estimate without assuming a fixed generative model?

Kernel density estimation (KDE) provides a solution.

- the approach is nonparametric: it keeps the entire training set
- we focus on the one-dimensional (continuous) case

High-level idea:

- Construct something similar to a histogram:
- For each data point, create a “bump” (via a Kernel)
- Sum up or average all the bumps



Kernel

KDE with a kernel $K: \mathbb{R} \rightarrow \mathbb{R}$:

$$p(x) = \frac{1}{n} \sum_{i=1}^n K(x - x_i)$$

K is adding bump around x_i , we need to keep all datapoints

e.g. slope of bump: $K(u) = \frac{1}{\sqrt{2\pi}} e^{-\frac{u^2}{2}}$, the standard Gaussian density.

Kernel needs to satisfy:

- symmetry: $K(u) = K(-u)$

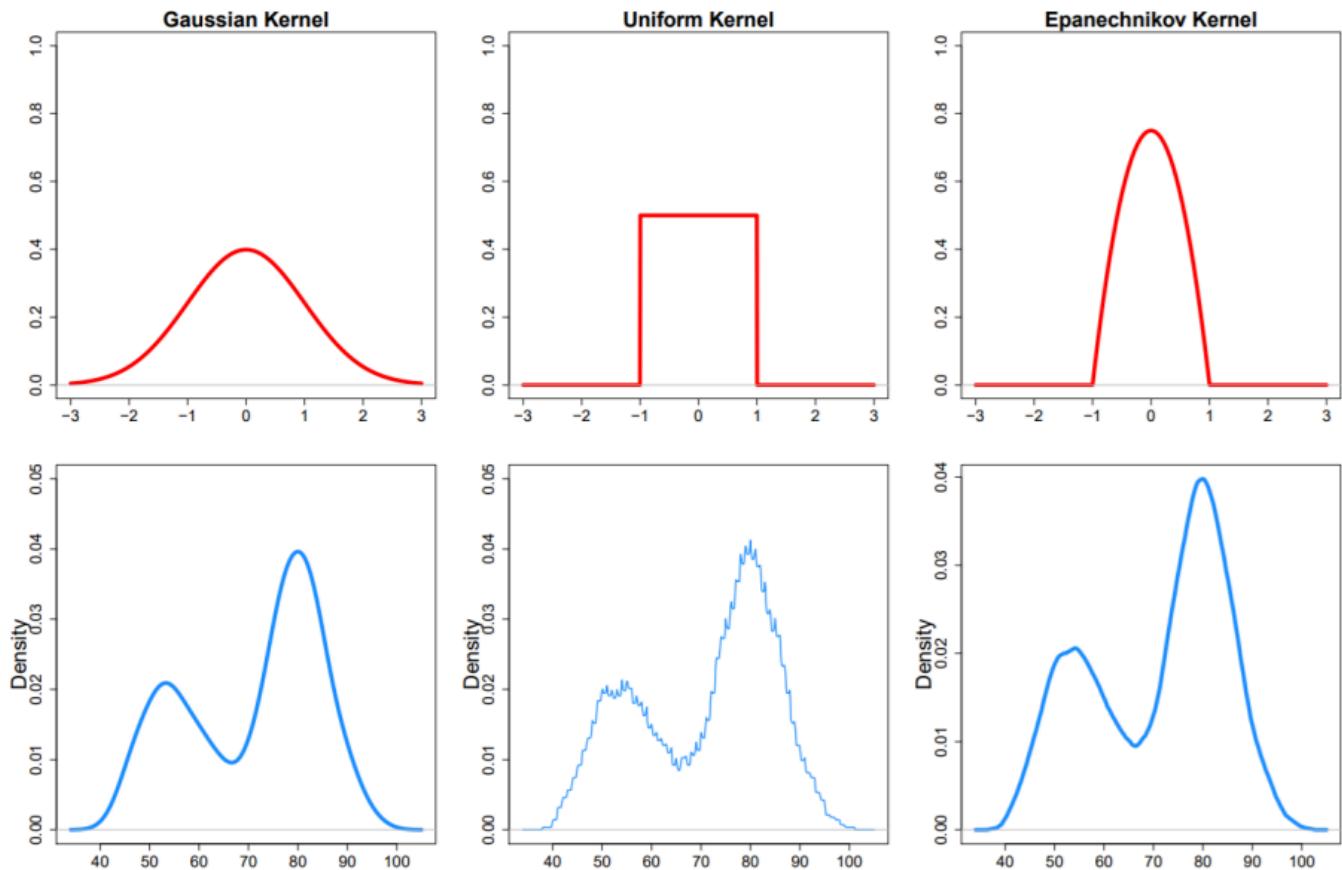
- $\int_{-\infty}^{\infty} K(u)du = 1$, makes sure p is a density function.

Different kernels

$$\frac{1}{\sqrt{2\pi}} e^{-\frac{u^2}{2}}$$

$$\frac{1}{2} \mathbb{I}[|u| \leq 1]$$

$$\frac{3}{4} \max\{1 - u^2, 0\}$$



Bandwidth

If $K(u)$ is a kernel, then for any $h > 0$

$$K_{h(u)} = \frac{1}{h} K\left(\frac{u}{h}\right) \quad (\text{stretching the kernel})$$

can be used as a kernel too (verify the two properties yourself) $\int_{-\infty}^{\infty} = \frac{1}{h} K\left(\frac{u}{h}\right) du = 1$.

if h increases,



So general KDE is determined by both the kernel K and the bandwidth h :

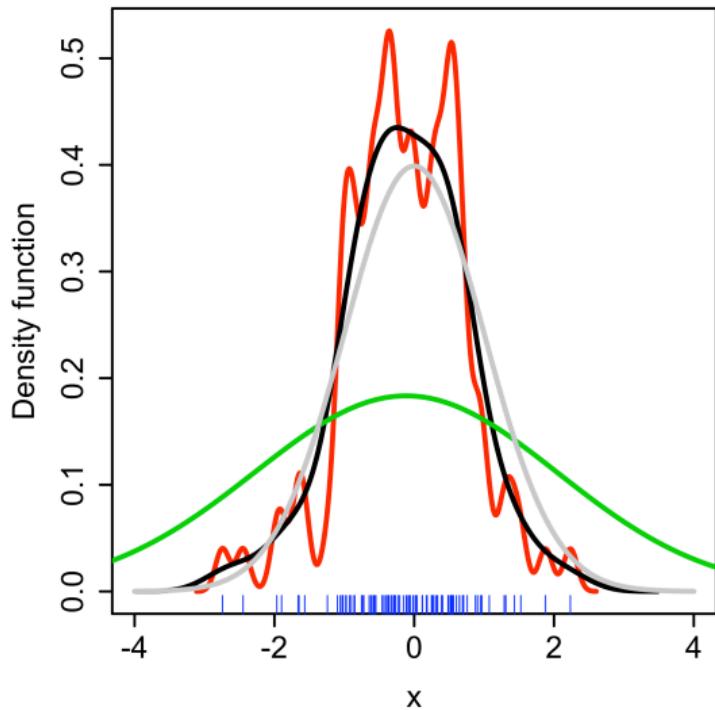
$$p(x) = \frac{1}{n} \sum_{i=1}^n K_h(x - x_i) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)$$

- x_i controls the center of each bump
- h controls the width/variance of the bumps.

Larger h means larger variance and smoother density.

Gray curve is ground-truth

- Red: $h = 0.05$
- Black: $h = 0.337$
- Green: $h = 2$



Naive Bayes

Suppose (x, y) is drawn from a joint distribution p . The Bayes optimal classifier is

$$f^*(x) = \arg \max_{c \in [C]} p(c|x)$$

i.e. predict the class with the largest conditional probability.

p is of course unknown, but we can estimate it, which is exactly a density estimation problem!

Estimation

How to estimate a joint distribution? Observe we always have

$$p(x, y) = p(y)p(x|y)$$

We know how to estimate $p(y)$ (C possible value) by now.

To estimate $p(x|y = c)$ for some $c \in [C]$, we are doing density estimation using data $\{x_i : y_i = c\}$.

This is not a 1D problem in general.

A naive assumption

Naive Bayes assumption: conditioning on a label, features are independent, which means

$$p(x|y = c) = \prod_{j=1}^d p(x_j|y = c)$$

x_j is the j -th coordinate of c . Now for each j and c we have a simple 1D density estimation problem!

Is this a reasonable assumption? Sometimes yes, e.g.

- use $x = (\text{Height}, \text{Vocabulary})$ to predict $y = \text{Age}$
- Height and Vocabulary are dependent
- but conditioned on Age, they are independent!

More often this assumption is unrealistic and "naive", but still Naive Bayes can work very well even if the assumption is wrong.

Example: Discrete features

Height: $\leq 3', 3' - 4', 4' - 5', 5' - 6', \geq 6'$ (5 possible values)

Vocabulary: $\leq 5K, 5K - 10K, 10K - 15K, 15K - 20K, \geq 20K$ (5 values)

Age: $\leq 5, 5 - 10, 10 - 15, 15 - 20, 20 - 25, \geq 25$ (6 possible values)

MLE estimation: e.g.

$$\begin{aligned} p(\text{Age} = 10 - 15) &= \frac{\#\text{examples with age } 10 - 15}{\#\text{examples}} \\ p(\text{Height} = 5' - 6' | \text{Age} = 10 - 15) \\ &= \frac{\#\text{examples with height } 5' - 6' \text{ and age } 10 - 15}{\#\text{examples with age } 10 - 15} \quad (\text{conditional} = \frac{\text{joint}}{\text{marginal}}) \end{aligned}$$

Discrete features: More formally

For a label $c \in [C]$

$$p(y = c) = \frac{|\{i : y_i = c\}|}{n}$$

For each possible value l of a discrete feature j ,

$$p(x_i = l | y = c) = \frac{|\{i : x_{i,j} = l, y_i = c\}|}{|\{i : y_i = c\}|}$$

$p(x_i = l | y = c)$ is the density for the j -th feature, $x_{i,j}$ means the j -th feature of i -th datapoint.

Continuous features

If the feature is continuous, we can do

- parametric estimation, e.g. via a Gaussian

$$p(x_j = x | y = c) = \frac{1}{\sqrt{2\pi}\sigma_{c,j}} \exp\left(-\frac{(x - \mu_{c,j})^2}{2\sigma_{c,j}^2}\right)$$

Where $\mu_{c,j}$ and $\sigma_{c,j}^2$ are the empirical mean and variance of feature j among all examples with label c .

- or nonparametric estimation, e.g. via a Kernel K and bandwidth h :

$$p(x_j = x | y = c) = \frac{1}{|\{i : y_i = c\}|} \sum_{i:y_i=c} K_h(x - x_{i,j})$$

Naive Bayes: Prediction

After learning the model

$$p(x, y) = p(y) \prod_{j=1}^d p(x_j | y)$$

the prediction for a new example x is

$$\begin{aligned} \arg \max_{c \in [C]} p(y = c | x) &= \arg \max_{c \in [C]} \frac{p(x, y = c)}{p(x)} = \arg \max_{c \in [C]} p(y = c) \cdot p(x | y = c) \\ &= \arg \max_{c \in [C]} \left(p(y = c) \prod_{j=1}^d p(x_j | y = c) \right) \quad (\text{Naive Bayes assumption, priors} \cdot \text{lielihood}) \\ &= \arg \max_{c \in [C]} \left(\ln p(y = c) + \sum_{j=1}^d \ln p(x_j | y = c) \right) \end{aligned}$$

For discrete features, plugging in previous MLE estimation gives

$$\begin{aligned} \arg \max_{c \in [C]} p(y = c | x) &= \arg \max_{c \in [C]} p(x, y = c) \\ &= \arg \max_{c \in [C]} \left(\ln p(y = c) + \sum_{j=1}^d \ln p(x_j | y = c) \right) \\ &= \arg \max_{c \in [C]} \left(\ln |\{i : y_i = c\}| + \sum_{j=1}^d \ln \frac{|\{i : x_{i,j} = x_j, y_i = c\}|}{|\{i : y_i = c\}|} \right) \end{aligned}$$

WHAT IS NAIVE BAYES LEARNING? CONNECTION TO LOGISTIC REGRESSION

Observe again for the case of continuous features with a Gaussian model, if we fix the variance for each feature to be σ (i.e. not a parameter of the model any more), then the prediction becomes

$$\begin{aligned}
 & \arg \max_{c \in [C]} p(y = c | x) \\
 &= \arg \max_{c \in [C]} \left(\ln |\{i : y_i = c\}| - \sum_{j=1}^d \left(\ln \sigma + \frac{(x_j - \mu_{c,j})^2}{2\sigma^2} \right) \right) \quad (x_j^2 \text{ is constant across all classes}) \\
 &= \arg \max_{c \in [C]} \left(\ln |\{i : y_i = c\}| - \sum_{j=1}^d \frac{\mu_{c,j}^2}{2\sigma^2} + \sum_{j=1}^d \frac{\mu_{c,j}}{\sigma^2} x_j \right) \\
 &= \arg \max_{c \in [C]} \left(w_{c0} + \sum_{j=1}^d w_{cd} x_j \right) = \arg \max_{c \in [C]} w_c^T x \quad (\text{linear classifier})
 \end{aligned}$$

where we denote $w_{c0} = \ln |\{i : y_i = c\}| - \sum_{j=1}^d \frac{\mu_{c,j}^2}{2\sigma^2}$ and $w_{cd} = \frac{\mu_{c,j}}{\sigma^2}$.

Moreover, by a similar calculation you can verify

$$p(y = c | x) \propto e^{w_c^T x}$$

This is the softmax function, the same model we used for the probabilistic interpretation of logistic regression!

So what is different then? They learn the parameters in different ways:

- both via MLE, logistic is on $p(y = c | x)$ (we said $p(y = c | x)$ for $x = \{\pm 1\} = \sigma(yw^T x)$), the other(Bayes) is on $p(x, y)$.
- solutions are different: logistic regression has no closed-form, naive Bayes admits a simple closed-form. (GMM doesn't have closed-form, either)

	Logistic regression	Naive Bayes
Model	conditional $p(y x)$	joint $p(x, y)$
Learning	MLE (can also be viewed as minimizing logistic loss)	MLE
Accuracy	usually better for large n	usually better for small n

Discriminative model

Generative model

Logistic Regression is a discriminative model (learn the classifier hypo-plane for each class), Naive Bayes is a generative model (model the density for each class).

MULTI-ARMED BANDITS (REINFORCEMENT LEARNING)

Decision making

Problems we have discussed so far:

- start with a fixed training dataset
- learn a predictor from the data or discover some patterns in the data

But many real-life problems are about learning continuously:

- make a prediction/decision
- receive some feedback
- repeat

Broadly, these are called online decision making problems.

Amazon/Netflix/MSN recommendation systems:

- a user visits the website
- the system recommends some products/movies/news stories
- the system observes whether the user clicks on the recommendation

Playing games (Go/Atari/StarCraft/...) or controlling robots:

- make a move
- receive some reward (e.g. score a point) or loss (e.g. fall down)
- make another move...

MULTIARMED BANDITS: MOTIVATION AND APPLICATIONS

Imagine going to a casino to play a slot machine

- it robs you, like a "bandit" with a single arm

Of course there are many slot machines in the casino

- like a bandit with multiple arms (hence the name)
- if I can play 10 times, which machines should I play?

This simple model and its variants capture many real-life applications:

- recommendation systems, each product/movie/news story is an arm (Microsoft MSN employs a variant of bandit algorithm)
- game playing, each possible move is an arm (AlphaGo has a bandit algorithm as one of the components)

Formal Setup

There are K arms (actions/choices/...)

The problem proceeds in rounds between the environment and a learner:

for each time $t = 1, \dots, T$

- the environment decides the reward for each arm $r_{t,1}, \dots, r_{t,K}$ (reward for each arm at time t).
- the learner picks an arm $\alpha_t \in [K]$

– the learner observes the reward for arm at time t

- the learner observes the reward for arm a_t , i.e., r_{t,a_t} .

Importantly, learner does not observe rewards for arms not selected!

This kind of limited feedback is usually referred to as bandit feedback.

Evaluating performance

What should be the goal here?

Maximizing total rewards $\sum_{t=1}^T r_{t,a_t}$ seems natural.

But the absolute value of rewards is not meaningful, instead we should compare it to some benchmark. A classic benchmark is

$$\max_{a \in [K]} \sum_{t=1}^T r_{t,a}$$

i.e. the largest reward one can achieve by always playing a fixed arm.

So we want to minimize

$$\max_{a \in [K]} \sum_{t=1}^T r_{t,a} - \sum_{t=1}^T r_{t,a_t}$$

This is called the regret: how much I regret not sticking with the best fixed arm in hindsight?

Environments

How are the rewards generated by the environments?

- they could be generated via some fixed distribution
- they could be generated via some changing distribution
- they could be generated even completely arbitrarily/adversarially

We focus on a simple setting:

- rewards($\{0, 1\}$) of arm a are i.i.d. samples of $Ber(\mu_a)$, that is, $r_{t,a}$ is 1 with prob. μ_a , and 0 with prob. $1 - \mu_a$, independent of anything else.
- each arm has a different mean (μ_1, \dots, μ_K) ; the problem is essentially about finding the best arm $\arg \max_a \mu_a$ as quickly as possible

Empirical means

Let $\hat{\mu}_{t,a}$ be the empirical mean of arm a up to time t :

$$\hat{\mu}_{t,a} = \frac{1}{n_{t,a}} \sum_{\tau \leq t : a_\tau = a} r_{\tau,a}$$

Where

$$n_{t,a} = \sum_{\tau \leq t} \mathbb{I}[a_\tau == a]$$

is the number of times we have picked arm a

is the number of times we have picked arm a .

Concentration: $\hat{\mu}_{t,a}$ should be close to μ_a if $n_{t,a}$ is large.

Exploitation only

Greedy: Pick each arm once for the first K rounds.

For $t = K + 1, \dots, T$, pick $a_t = \arg \max_a \hat{\mu}_{t-1,a}$

What's wrong with this greedy algorithm?

Consider the following example:

- $K = 2, a_1 \sim \text{Bern}(\mu_1 = 0.6), a_2 \sim \text{Bern}(\mu_2 = 0.5)$ (so arm 1 is the best)
- suppose the algorithm first picks arm 1 and sees reward 0, then picks arm 2 and sees reward 1 (this happens with decent probability: $0.4 \times 0.5 = 0.2$)
- the algorithm will never pick arm 1 again, opposed to optimal solution.

All bandit problems face the same dilemma: **Exploitation vs. Exploration trade-off**

- on one hand we want to exploit the arms that we think are good.
- on the other hand we need to explore all arms often enough in order to figure out which one is better.
- so each time we need to ask: do I explore or exploit? and how?

We next discuss three ways to trade off exploration and exploitation for our simple multi-armed bandit setting.

Explore-then-Exploit

A natural first attempt: Explore-then-Exploit.

Input: a parameter $T_0 \in [T]$.

Exploration phase: for the first T_0 rounds, pick each arm for T_0/K times.

Exploitation phase: for the remaining $T - T_0$ rounds, stick with the empirically best arm $a_t = \arg \max_a \hat{\mu}_{T_0,a}$.

Parameter T_0 clearly controls the exploration/exploitation trade-off.

Issues: It's pretty reasonable, but the disadvantages are also clear:

- not clear how to tune the hyper-parameter T_0
- in the exploration phase, even if an arm is clearly worse than others based on a few pulls, it's still pulled T_0/K times
- clearly it won't work if the environment is changing

ϵ -Greedy: A slightly better algorithm

ϵ -Greedy Pick each arm once for the first K rounds.

For $t = K + 1, \dots, T$,

- with probability ϵ , explore: pick an arm uniformly at random
- with probability $1 - \epsilon$, exploit: pick $a_t = \arg \max_a \hat{\mu}_{t-1,a}$.

Pros

- always exploring and exploiting
- applicable to many other problems
- first thing to try usually

Cons

- need to tune ϵ
- same uniform exploration

Is there a more adaptive way to explore?

UCB: More adaptive exploration

A simple modification of "Greedy" leads to the well-known:

Upper Confidence Bound (*UCB*) algorithm

For $t = 1, \dots, T$, pick $a_t = \arg \max_a UCB_{t,a}$ where

$$UCB_{t,a} = \hat{\mu}_{t-1,a} + 2\sqrt{\frac{\ln t}{n_{t-1,a}}}$$

- $UCB_{t,a}$ initialize all estimates to be same.
- If $n_{t-1,a}$ is small, $UCB_{t,a}$ will be large.
- the first term in $UCB_{t,a}$ represents exploitation, while the second (bonus) term represents exploration.
- the second (bonus) term is large if the arm is not pulled often enough, which encourages exploration (adaptive due to the first term)
- a parameter-free algorithm, and it enjoys optimal regret!

Why is it called upper confidence bound?

One can prove that with high probability:

$$\mu_a \leq UCB_{t,a}$$

so $UCB_{t,a}$ is indeed an upper bound on the true mean.

Another way to interpret *UCB*, "optimism in face of uncertainty":

- true environment (best mean) is unknown due to randomness (uncertainty)
- have an upper bound (optimistic guess) on the expected reward of each environment, and pick best one according to upper bound (optimism)

This principle is useful for many other bandit problems.

Limitations of multi-armed bandits

Multi-armed bandit is among the simplest decision making problems with limited feedback.



Often, it can be too simple to capture real-life problems. One important aspect it fails to capture is the “state” of the learning agent, which has impacts on the reward of each action.

- e.g. for Atari games, after making one move, the agent moves to a different state, with possible different rewards for each action.

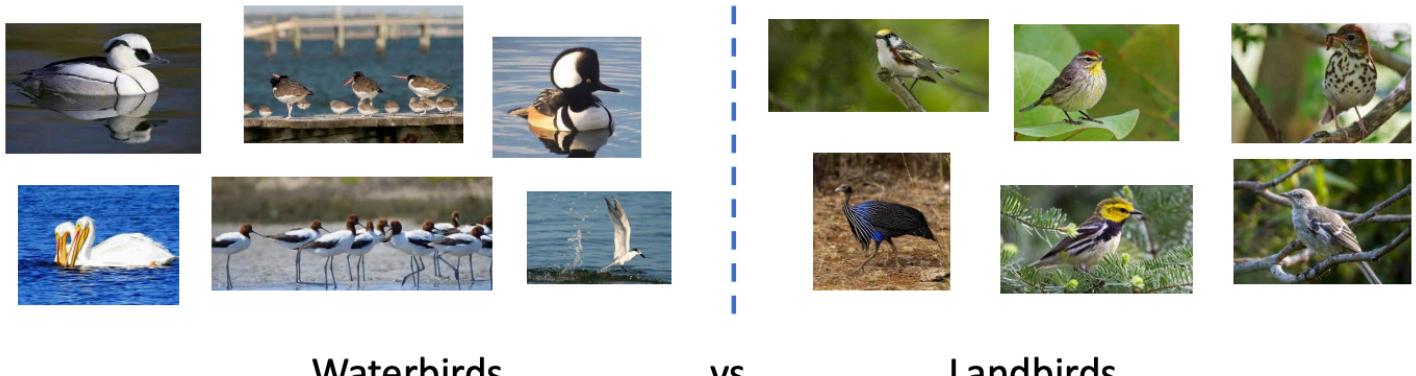
There are many other techniques and models in reinforcement learning (RL) which can deal with this issue.

Responsible ML

ML models can be very sensitive to changes in the data distribution

ML models can latch onto spurious features to make predictions.

You saw a small example of this in the HW3 Bonus question:



Waterbirds

vs.

Landbirds

Most images of waterbirds are in water, and landbirds are on land. But this isn't always true! This is known as failure to distributional shifts.





Waterbirds

vs.

Landbirds

A real-world example: CNN models have obtained impressive results for diagnosing X-rays. But the models may not generalize as well to data from new hospitals because they can learn to pickup on spurious correlations such as the type of scanner and marks used by technicians in specific hospitals!

A real-world example

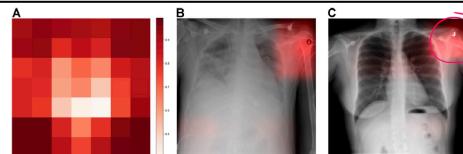
CNN models have obtained impressive results for diagnosing X-rays

E.g. *ChestX-ray8: Hospital-scale Chest X-ray Database and Benchmarks on Weakly-Supervised Classification and Localization of Common Thorax Diseases*, Wang et al., 2017



Source: *Deep learning for chest radiograph diagnosis: A retrospective comparison of the CheXNeXt algorithm to practicing radiologists*, Rajpurkar et al. 2018

But the models may not generalize as well to data from new hospitals because they can learn to pickup on spurious correlations such as the type of scanner and marks used by technicians in specific hospitals!



CNN to predict hospital system detects both general and specific image features.
(A) We obtained activation heatmaps from our trained model and averaged over a sample of images to reveal which subregions tended to contribute to a hospital system classification decision. Many different subregions strongly predicted the correct hospital system, with especially strong contributions from image corners. (B-C) On individual images, which have been normalized to highlight only the most influential regions and not all those that contributed to a positive classification, we note that the CNN has learned to detect a metal token that radiology technicians place on the patient in the corner of the image field of view at the time they capture the image. When these strong features are correlated with disease prevalence, models can leverage them to indirectly predict disease.

Source: *Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: A cross-sectional study*, Zech et al. 2018

ML models can be biased against certain sub-populations

You saw a small example of this in the HW4 word embedding question.

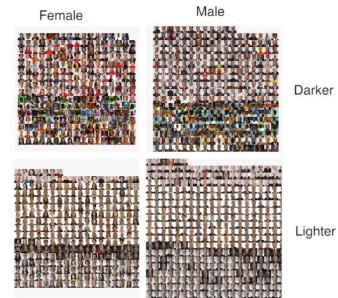
The Gendershades Project. How well do facial recognition tools perform on various demographics? Ans: Not very well.





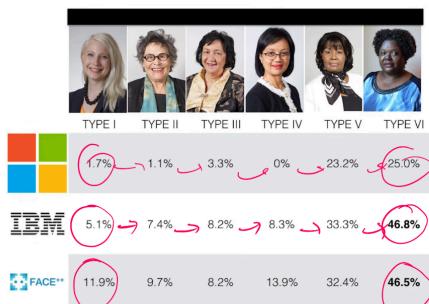
The Gendershades Project

How well do facial recognition tools perform on various demographics?



<http://gendershades.org/>

Ans: Not very well



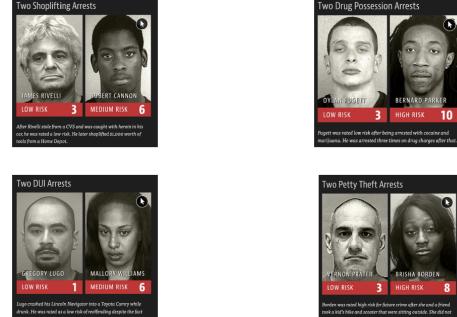
Ans: Not very well



The COMPAS software.

The COMPAS software

There's software used across the country to predict future criminals. And it's biased against blacks.



<https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>

The New York Times

There Is a Racial Divide in Speech-Recognition Systems, Researchers Say

Technology from Amazon, Apple, Google, IBM and Microsoft misidentified 35 percent of words from people who were black. White people fared much better.



[Link to article](#)

MIT Technology Review

ARTIFICIAL INTELLIGENCE

LinkedIn's job-matching AI was biased. The company's solution? More AI.

ZapRecruiter, CoverBuilder, LinkedIn—most of the world's biggest job search sites use AI to match people with job openings. But the algorithms don't always play fair.

By Sheridan Hall & Mike Schelmann June 23, 2021

[Link to article](#)

How an AI grading system ignited a national controversy in the U.K.

Bryan Head, author of *AI Class*



A huge controversy in the U.K. over an algorithm used to substitute for university-entrance exams highlights problems with the use of AI in the real world.

[Link to article](#)

The Atlantic

TECHNOLOGY

It Was Supposed to Detect Fraud. It Wrongfully Accused Thousands Instead.

How Michigan's attempt to automate its unemployment system went horribly wrong

By Stephanie Wykstra and Undark



[Link to article](#)

Conclusion

This class:

- Understand the fundamentals
- Understand when ML works, its limitations, think critically

In particular,

- Study fundamental statistical ML methods (supervised learning, unsupervised learning, etc.)
- Solidify your knowledge with hand-on programming tasks
- Prepare you for studying advanced machine learning techniques

1. Examine your data
2. Examine your model

ML/AI can be very powerful, but should be used responsibly.

Lots of great resources to learn more about best AI/ML practices.

<https://ai.google/responsibilities/responsible-aipractices/>